





Capabilities Development: 3
Easy Pieces

Today

- ARM
- Some Syscalls
- File IO
- Tour of the lab

What is ARM?

- stands for **Advanced RISC Machine(s)**.
- RISC stands for reduced instruction set computer
- ARM is a company, an ambiguous term for Assembly Languages and architectures

ARM Holdings

- A British(?) semiconductor and software company.
- Specializes in designing RISC architectures.
- Does not manufacture its own chips.
- Licenses its designs to other companies (e.g., Qualcomm, Apple, Samsung).
- ARM chips power billions of devices, from smartphones to IoT devices to SSD controllers to enterprise servers.
- Majority owned by Soft Bank (questionable call by the UK)

RISC vs CISC

Feature	RISC	CISC
Instruction Set	Simple, small	Complex, large
Execution	Single cycle per instruction	Multi-cycle per instruction
Power Efficiency	High	Lower
Performance	Optimized for pipelining	Slower due to complexity
Program Size	Larger	Smaller
Examples	ARM, RISC-V	x86, Intel 8086

(Note x86 is kind of like a sugared RISC at this point)

ArmV8 profiles

1. **A (Application):**

- Supports rich operating systems.
- Focused on performance and complex applications.
- Examples: smartphones, tablets, servers.

2. **R (Real-Time):**

- Predictable and deterministic performance.
- Common in automotive and industrial systems.

3. **M (Microcontroller):**

- Low power, low cost.
- Designed for IoT and embedded systems.
- probably in your SSD controller

Microarchitectures

- Specific implementations of the ARM architecture.
- Examples: Cortex-A series, Cortex-R series, Cortex-M series.
- Each microarchitecture optimizes for specific use cases:
 - Performance, power efficiency, or a balance of both.

ARMv8-a

- This class focuses on ARMv8-a
- RISC arch built by ARM holdings
- aarch64 -> A64: 64 bit execution state
- aarch32 -> A32: 32 bit execution state

A64 (AArch64 mode) Assembly Crash Course

- aarch64: 64 bit execution state for ARMv8-a
- A64 Assembly language/instruction set for aarch64
- 32-bit assembly instructions (4 byte word)
- Uses 31 64 bit general-purpose registers: x0–x30, plus sp (stack pointer) and pc (program counter).
 - w0–w30 are the lower 32 bits b_0, \dots, b_{31}
- Uses 31 128 bit floating point registers q0–q31
 - b, h, w, d, q: ->
 - byte, half-word, word, double-word, quad-word
- Some registers have special usage/convention:
- The goal here is to learn **just enough** to be useful
- technically supports little/big endian but I have only ever seen little endian

Special General Purpose registers (A64)

- Here "special" means there exists a convention
- x0 often used for function return values from a subroutine.
- x0–x7 typically used for arguments in many Linux ABIs.
- x8 holds syscall number
- x29: (sometimes aliased as fp) Frame Pointer (optional)
- x30: Link Register often holds return address on subroutine calls

Comparison to x86_64

Aarch64 Register	Purpose/Usage	x86_64 Equivalent	Explanation
x0	Return values from a subroutine	rax	Both are used to store return values for functions.
x0-x7	Arguments in many Linux ABIs	rdi, rsi, rdx, rcx, r8, r9	x86_64 uses a similar approach with a set of registers for function arguments.
x8	Syscall number	rax	On x86_64, the syscall number is passed through rax before invoking syscall.
x29 (aliased fp)	Frame pointer (optional; helps manage stack frames)	rbp	Both are used as a frame pointer to access local variables and manage the stack.
x30	Link register (holds return address for subroutine calls)	ret stack mechanism	In x86_64, the return address is stored on the stack instead of in a dedicated register.

A64 vs x86_64 notes

- Fixed-length 32-bit instructions simplify decoding and pipelining.
- Implements a **load/store architecture**, meaning it cannot perform operations directly on memory. - Aarch64 cannot directly operate on data unless it is stored in a register.
- All data must first be loaded into registers for processing and then stored back into memory.
 - For example, in memcpy, data is copied in chunks by loading blocks into registers and then storing them to the destination address.

Other important registers

- xzr: zero register. ignores writes.
- sp: stack pointer
- pc: instruction pointer/program counter. Can't directly read/modify like x_i
 - bracnh: used to modify pc
- lr: alias for x30.
 - There is no call instruction.
 - There is no ret instruction
- various system registers (exception levels, mmu ...etc)
- this is more trivia than anything but "x31" is either zero register or stack pointer
 - sp and xzr/wzr are **architecturally distinct**.
 - in *some instruction encodings* the same field in the instruction word can represent either sp or xzr depending on context.

Refresher: Common Registers (AArch64)

- **x0–x7**:
- **x8**:
- **x9–x15**: temporary registers.
- **x16–x17**: intra-procedure call scratch regs.
- **x19–x28**: callee-saved registers.
- **x29**:
- **x30**:
- **sp**: stack pointer.
- **pc**: program counter (auto-updated cant be directly used).

Refresher: Common Registers (AArch64)

- **x0–x7**:
 - used to pass function arguments, return values.
- **x8**:
- **x9–x15**: temporary registers.
- **x16–x17**: intra-procedure call scratch regs.
- **x19–x28**: callee-saved registers.
- **x29**:
- **x30**:
- **sp**: stack pointer.
- **pc**: program counter (auto-updated cant be directly used).

Refresher: Common Registers (AArch64)

- **x0-x7**:
 - used to pass function arguments, return values.
- **x8**:
 - indirect syscall param or special usage.
- **x9-x15**: temporary registers.
- **x16-x17**: intra-procedure call scratch regs.
- **x19-x28**: callee-saved registers.
- **x29**:
- **x30**:
- **sp**: stack pointer.
- **pc**: program counter (auto-updated cant be directly used).

Refresher: Common Registers (AArch64)


- **x0–x7**:
 - used to pass function arguments, return values.
- **x8**:
 - indirect syscall param or special usage.
- **x9–x15**: temporary registers.
- **x16–x17**: intra-procedure call scratch regs.
- **x19–x28**: callee-saved registers.
- **x29**:
 - optional frame pointer (often used with fp alias)
- **x30**:
- **sp**: stack pointer.
- **pc**: program counter (auto-updated cant be directly used).

Refresher: Common Registers (AArch64)

- **x0–x7**:
 - used to pass function arguments, return values.
- **x8**:
 - indirect syscall param or special usage.
- **x9–x15**: temporary registers.
- **x16–x17**: intra-procedure call scratch regs.
- **x19–x28**: callee-saved registers.
- **x29**:
 - optional frame pointer (often used with fp alias)
- **x30**:
 - link register (return address).
- **sp**: stack pointer.
- **pc**: program counter (auto-updated cant be directly used).

Operand Sizes/ Types

- Byte: 8 bits
- Halfword: 16 bits
- Word: 32 bits
- Doubleword: 64 bits
- Quadword: 128 bits
- ARMv8-a is (almost always) **little-endian** *.
- In assembly, you'll see mnemonics for loading/storing different widths (e.g. `ldrb`, `ldrh`, `ldr`, `ldur`).



Sanity Check

What does this print? #exercise

```
#include <stdio.h>

int main() {
    int x = 0xdeadbeef;
    unsigned char *y = (unsigned char *)&x;
    for (int i = 0; i < 4; i++) {
        unsigned char c = y[i] & 0xff;
        printf("%x ", c);
    }
    printf("\n");
}
```

Sanity Check

What does this print? #exercise

```
#include <stdio.h>

int main() {
    int x = 0xdeadbeef;
    unsigned char *y = (unsigned char *)&x;
    for (int i = 0; i < 4; i++) {
        unsigned char c = y[i] & 0xff;
        printf("%x ", c);
    }
    printf("\n");
}
```

- ef be ad de

Basic Arithmetic (AArch64)

- **ADD / SUB**: Integer addition/subtraction
 - Example: `ADD x0, x1, x2` ($x0 \leftarrow x1 + x2$)
 - Immediate form: `ADD w3, w4, #10`
- **ADDS / SUBS**: Same as above, but **sets condition flags** (N, Z, C, V)
 - Useful for branching on results
- **MUL**: Multiply the lower 64 bits
 - `MUL x0, x1, x2`
- **SMULH / UMLH**: Signed / unsigned high 64-bit multiply
 - If the product exceeds 64 bits, high part stored in the destination

Load/Store Operations

- **LDR / STR**: Primary load/store instructions
 - `LDR x0, [x1] → x0 ← (uint64_t)x1`
 - `STR x2, [x3, #16] → (uint64_t)(x3+16) ← x2`
- **LDRB / STRB**: For 8-bit (byte) load/store
 - Similarly LDRH, LDRW for halfword/word
- **LDP / STP**: Load/Store **pairs** of registers
 - Often used to save/restore register pairs in function prolog/epilog
- Offsets can be
 - immediate (`[x1, #4]`)
 - post-/pre-indexed (`[x1], #4, [x1, #4]!`)
- READ: <https://developer.arm.com/documentation/102374/0102/Loads-and-stores---addressing>

Immediate Offset

- The address is computed by adding a constant offset directly to the base register.
- the base register is not modified.

Syntax:

```
ldr w0, [x1, #4] // Load from address (x1 + 4)
str w0, [x1, #-8] // Store to address (x1 - 8)
```

- address = $x1 + 4$ (or $x1 - 8$).
- Perform the memory access.
- Base register x1 remains unchanged.

Post-indexed Offset

- The address is computed using the base register.
- After the access, the offset is added to the base register.

Syntax:

```
ldr w0, [x1], #4    // Load from x1, then x1 += 4  
str w0, [x1], #-8   // Store to x1, then x1 -= 8
```

- Use the address in the base register for memory access.
- Update the base register ($x1 += 4$ or $x1 -= 8$).

Pre-Index Offset

- The offset is added to the base register before the memory access.
- The base register is updated with the new address.

Syntax:

```
ldr w0, [x1, #4]! // x1 += 4, then load from x1  
str w0, [x1, #-8]! // x1 -= 8, then store to x1
```

- Update the base register ($x1 += 4$ or $x1 -= 8$).
- Use the updated value of the base register for memory access.

Conditional Flags

- Condition Codes: Special flags set by the processor to indicate the result of an operation.
- Enables conditional execution of instructions based on previous operations.
- N (Negative): Set if the result of an operation is negative.
- Z (Zero): Set if the result is zero.
- C (Carry): Set if an operation results in a carry out or borrow.
- V (Overflow): Set if an operation causes a signed overflow.

Instructions-> Conditional Flags

- Comparison Instructions:
 - CMP: Compares two values by subtracting one from the other; updates N, Z, C, V flags.
 - CMN: Compares two values by adding them; updates N, Z, C, V flags.
- Logical Instructions:
 - TST: Performs an AND operation; updates N and Z flags.
 - TEQ: Performs an XOR operation; updates N and Z flags.

Basic Control Flow


- **B** : Unconditional branch
- **BL** : Branch and link (subroutine call)
 - Automatically saves the return address in **x30** (Link Register)
- **RET**: Return from subroutine
 - Not a real instruction :)
 - `blr x30`
 - Jumps to the address in x30
- **CBZ / CBNZ**: Compare register to zero and branch
 - `CBZ x0, label` → if `x0 == 0`, branch
- Condition-based branches (after `ADDS`, `SUBS`):
 - **B.EQ**, **B.NE**, **B.GT**, **B.LT**, etc.

More Resources

- <https://developer.arm.com/documentation/102374/0102>
- <https://web.archive.org/web/20240829145252/https://modexp.wordpress.com/2018/10/30/arm64-assembly/>
- <https://mariokartwii.com/armv8/>

Reading the Docs

- “RTFM” (read the friendly manual) is vital for learning about Linux.
- Example: `man 2 open`, `man 3 printf`.
- to learn about the man pages,
 - `$ man man`



RTFM

- Debugging your code for 8 hours can save you 5 minutes of reading the docs
 - I myself, routinely don't read the documentation and suffer for it. Be better than me. Learn from my mistakes. RTFM

1337 Documentation reader



wrote the top
answer on stackoverflow

script kiddies



stole code from the
top answer on stackoverflow

How do we learn A64 in this class?

- `objdump -d a.out`
- Writing little bootstraps

Baby's first A64

```
.section .text
.global _start

_start:
    mov x8, #93          // Syscall number for exit (93)
    mov x0, #0           // Exit code (0)
    svc #0               // Make the syscall
```

Baby's first A64

```
.equ sys_exit, 93

.section .text
.global _start

_start:
    mov x8, sys_exit          // Syscall number for exit (93)
    mov x0, #0                // Exit code (0)
    svc #0                    // Make the syscall
```

Baby's first A64

```
// Define sys_exit to be 93
.equ sys_exit, 93

// Macro to perform the exit syscall
// note that code is an arg
.macro exit code
    mov x8, sys_exit    // Load syscall number (94 for exit)
    mov x0, \code       // Load the exit code into x0
    svc #0              // Make the syscall
.endm

.section .text
.global _start

_start:
    exit 0              // Call the macro with exit code 0
```

Baby's first A64

```
// common.S
// Define syscall numbers
.equ sys_exit, 93 // Syscall number for exit

// Macro to perform the exit syscall
.macro exit code
    mov x8, sys_exit    // Load syscall number
    mov x0, \code        // Load the exit code into x0
    svc #0              // Make the syscall
.endm
```

```
.include "common.S" // Include the macros file

.section .text
.global _start

_start:
    exit 0
```

hello world

- write: fd=1, addr_str, str_size

```
.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
_start:
    // 1 is stdout
    mov x0, #1
    ldr x1, =hello
    mov x2, #14
    mov x8, #64
    svc 0
    mov x0, 0
    mov x8, #93
    svc 0
```


hello world

```
// Macros
.equ STDOUT_FD, 1
.equ SYS_WRITE, 64
.equ  SYS_EXIT, 93

.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
_start:
    // 1 is stdout
    mov x0, STDOUT_FD
    ldr x1, =hello
    mov x2, #14
    mov x8, SYS_WRITE
    svc 0
    mov x0, 0
    mov x8, SYS_EXIT
    svc 0
```

hello world

```
// Macros
.equ STDOUT_FD, 1
.equ SYS_WRITE, 64
.equ SYS_EXIT, 93

.macro write_stdout message,length
    mov x0, STDOUT_FD
    ldr x1, =\message
    mov x2, \length
    mov x8, SYS_WRITE
    svc 0
.endm

.macro exit_with code
    mov x0, \code
    mov x8, SYS_EXIT
    svc 0
.endm

.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
start:
```

hello world

```
.equ  SYS_EXIT, 93
.section .rodata
fmt: .asciz "Hello %s!\n"
fmt_alt: .asciz "fmt addr: 0x%p\n"
val: .asciz "world"

.section .text
.global _start

_start:
    ldr x0, =fmt
    ldr x1, =val
    bl printf
    ldr x0, =fmt_alt
    ldr x1, =fmt
    bl printf
    mov x8, SYS_EXIT
    svc 0
```

- <https://godbolt.org/z/Gxjor9Kqj>

```
#include <stdio.h>


unsigned long long factorial(unsigned long long n){
    if (n == 0){
        return 1;
    }
    return n * factorial(n-1);
}

int main(int argc, char** argv){
    unsigned long long out = factorial(5);
    printf("Out: %llu\n", out);
    return 0;
}
```

```
factorial(unsigned long long):                                // @factorial(unsigned long long)
    cbz     x0, .LBB0_4
    mov     x8, x0
    mov     w0, #1
.LBB0_2:                                                    // =>This Inner Loop Header: Depth=1
    subs    x9, x8, #1
    mul     x0, x8, x0
    mov     x8, x9
    b.ne    .LBB0_2
    ret
.LBB0_4:
    mov     w0, #1
    ret
main:                                                    // @main
    str     x30, [sp, #-16]!                                // 8-byte Folded Spill
    adrp    x0, .L.str
    add     x0, x0, :lo12:.L.str
    mov     w1, #120
    bl      printf
    mov     w0, wzr
    ldr     x30, [sp], #16                                  // 8-byte Folded Reload
    ret
.L.str:
    .asciz  "Out: %llu\n"
```

Linux I/O on AArch64: Syscalls


- Focus: file I/O syscalls on AArch64
- I.e. "How do we read and filter through files"
- Use direct `syscall()`
- Explore kernel internals & structs
- Build efficient scanner for a 2GB file
- In-class assignment: structured raw file scan



Finding Syscall values

```
cat /usr/include/asm-generic/unistd.h | grep openat  
#define __NR_openat 56  
__SYSCALL(__NR_openat, sys_openat)  
#define __NR_openat2 437  
__SYSCALL(__NR_openat2, sys_openat2)
```

```
grep -ho "__NR_[a-zA-Z0-9_]\+\s\+[0-9]\+" /usr/include/asm-generic/unistd.h | \  
sed 's/__NR_/' | column -t
```



Syscall Interface (AArch64)

- Syscalls invoked with `svc #0`
- Registers:
 - `x8` = syscall number
 - `x0–x5` = up to 6 args
 - return value in `x0`
- Example:

```
int fd = syscall(SYS_openat, AT_FDCWD, "file.txt", O_RDONLY);
```

Kernel Objects Overview

```
task_struct
├── files → files_struct
│   └── fd table → struct file *
├── mm → mm_struct
│   └── vm_area_struct list/tree
```

Kernel Objects

- How to find kernel structs in linux
- `struct file` — <https://github.com/torvalds/linux/blob/4ff71af020ae59ae2d83b174646fc2ad9fcd4dc4/include/linux/fs.h#L1099>
- `struct file_operations` <https://github.com/torvalds/linux/blob/4ff71af020ae59ae2d83b174646fc2ad9fcd4dc4/include/linux/fs.h>
- `struct mm_struct` [https://github.com/torvalds/linux/blob/4ff71af020ae59ae2d83b174646fc2ad9fcd4dc4/include/linux/mm_types.h](https://github.com/torvalds/linux/blob/4ff71af020ae59ae2d83b174646fc2ad9fcd4dc4/include/linux/mm_types.h)
- `struct vm_area_struct` https://github.com/torvalds/linux/blob/4ff71af020ae59ae2d83b174646fc2ad9fcd4dc4/include/linux/mm_types.h

man openat

- Open file relative to directory (or AT_FDCWD for cwd)
 - Or create depending on arguments O_*
- On success, kernel creates struct file, updates task_struct->files
 - Inserts into fd table

Userland:

```
#include <syscall.h>
...
int fd = syscall(SYS_openat, AT_FDCWD, "/tmp/ch0nky.txt", O_RDONLY);
```

man read

- Reads bytes into user buffer
- Kernel uses page cache, `copy_to_user(...)`
- Updates `file->f_pos`

Userland:

```
char buf[4096];  
ssize_t n = syscall(SYS_read, fd, buf, sizeof(buf));
```

man write

- Copies data from user → kernel
- Updates page cache, marks pages dirty
- Advances file → f_pos
- Logically used to send data to an object managed by the kernel (file, pipe,..etc)

Userland:

```
const char *msg = "Hello\n";
syscall(SYS_write, fd, msg, strlen(msg));
// example: writing data to stdout
syscall(SYS_write, 1, msg, strlen(msg));
```

stat / fstat

- Retrieves file metadata from inode
- No new file object created
- Useful for size, mode, timestamps

Userland:

```
struct stat st;  
syscall(SYS_fstat, fd, &st);  
printf("Size: %lld\n", (long long) st.st_size);
```

man lseek

- Moves file offset (file->f_pos)
- SEEK_SET, SEEK_CUR, SEEK_END
- Only for seekable fds

Userland:

```
off_t size = syscall(SYS_lseek, fd, 0, SEEK_END);
```


man close

- Releases fd from files_struct
- Decrements struct file refcount
- May free file object

Userland:

```
syscall(SYS_close, fd);
```

man pread / pwrite

- pread(fd, buf, count, offset)
- Reads from fd at offset
- Does not change file->f_pos
- Atomic (no race lseek+read)
- pwrite = write at offset
- Syscall: __NR_pread64 (AArch64 = 67)

Userland:

```
char buf[16];  
ssize_t n = syscall(SYS_pread64, fd, buf, 16, 100);
```

man mmap

- Maps file region into process memory
- Creates new vm_area_struct in mm_struct
- Pages loaded lazily on fault

Userland:

```
char *map = syscall(SYS_mmap, NULL, size,  
                    PROT_READ, MAP_PRIVATE, fd, 0);
```

man munmap

- munmap: removes VMA from mm_struct
- Kernel updates VMA flags + page tables

Userland:

```
syscall(SYS_munmap, map, size);
```

man mprotect

- mprotect: changes page protections
- Kernel updates VMA flags + page tables

Userland:

```
syscall(SYS_munmap, map, size);
```

Efficient Large File Scanning

Goal: find lines starting with "FLAG{" in 2GB file.

Steps:

- openat file
- fstat size
- mmap whole file (or chunked)
- Scan for prefix after newline/start
- munmap + close

Example Scanner

```
char *data = mmap(NULL, size, PROT_READ, MAP_PRIVATE, fd, 0);
for (off_t i = 0; i < size; i++) {
    if (i == 0 || data[i-1] == '\n') {
        if (memcmp(&data[i], "FLAG{", 5) == 0) {
            // Found line
        }
    }
}
munmap(data, size);
```

Discussion

- Implant developer's perspective: what uses of file IO might we need?